ACTSC432 — Loss Models II

Classnotes for Spring 2019

by

Johnson Ng

BMath (Hons), Pure Mathematics major, Actuarial Science Minor University of Waterloo

Table of Contents

Ta	ıble o	f Contents	2
Li	st of	Definitions	3
Li	st of	Theorems	4
Li	st of	Procedures	5
Pr	eface	•	7
1	Intr	oduction and Review of Probability	9
	1.1	Introduction to Credibility Theory	9
	1.2	Review of Probability	11
2	Rev	riew of Statistics	13
	2.1	Unbiased Estimation	13
	2.2	Mean Squared Error	18
	2.3	Maximum Likelihood Estimation	19
	2.4	Bayesian Estimation	21
		2.4.1 Conjugate Prior Distributions and the Linear Exponential Family	23
3	Lim	nited Fluctuation Credibility Theory	27
	3.1	Limited Fluctuation Credibility	27
	3.2	Full Credibility	28
	3.3	Partial Credibility	33
	3.4	Problems with Limited Fluctuation Credibility	36
Bi	bliog	graphy	37

Index 38

List of Definitions

1	■ Definition (Estimate)	13
2	■ Definition (Estimator)	13
3	■ Definition (Biased and Unbiased Estimator)	14
4	■ Definition (Sample Variance)	16
5	■ Definition (Mean Squared Error)	18
6	■ Definition (Likelihood Function)	19
7	■ Definition (Maximum Likelihood Estimation)	19
8	■ Definition (Log-likelihood Function)	19
9	■ Definition (Prior Distribution)	21
10	■ Definition (Joint Distribution)	22
11	■ Definition (Marginal Distribution)	22
12	■ Definition (Posterior Distribution)	22
13	■ Definition (Posterior Mean)	23
14	■ Definition (Bayes Estimator)	23
15	■ Definition (Conjugate Prior Distribution)	24
16	■ Definition (Linear Exponential Family)	24

List of Theorems

1	• Proposition (Sample Mean as the Unbiased Estimator of the Mean)	16
2	♦ Proposition (Sample Variance as the Unbiased Estimator of the Variance)	16
3	♦ Proposition (Formula for the Posterior Distribution)	23
4	■Theorem (Conjugate Prior Distributions of Linear Exponential Distributions)	25

List of Procedures



For this set of notes, I shall follow the format of which the course is presented, by breaking contents into modules instead of lectures. Also, I will be relying on the standard textbook for this topic, namely Klugman et al. 2012.

1 Introduction and Review of Probability

We shall first take an overview of what this course is about, and we will review on some of the relevant notions from earlier courses.

1.1 Introduction to Credibility Theory

Credibility Theory is a form of statistical inference that

- uses newly observed past events; to
- more accurately re-forecasts uncertain future events.

From Klugman et al. 2012,

It is a set of quantitative tools that allows an insurer to perform prospective experience rating (adjust future premiums based on past experience) on a risk or group of risks. If the experience of a policyholder is consistently better than that assumed in the underlying manual rate (also called a pure premium), then the policyholder may demand a rate reduction.

That's all fancy mumbo-jumbo so let's go through an example that will hopefully enlighten us.

Example 1.1.1 (Enlightening Example to Credibility Theory)

Suppose automobile insurance policies are classified according to the following factors:

• number of drivers;

- gender of each driver;
- number of vehicles; and
- brand, model, production year, and approximate mileage driver per year.

Policies with identical characteristics are assumed to belong to the same **rating class**, which represents a group of individuals with similar risks.

Suppose there are 2 policies in the same rating class. Both policies are charged with a so-called **manual premium** of \$1,500 per year. This is the premium specified in the insurance manual for a policy with similar characteristics.

Let's say that after 3 years, we obtain the following data: We want

	Policy 1	Policy 2
Year 1	0	500
Year 2	200	4000
Year 3	0	2500

to find out what's a good premium to charge to each policy for Year 4.

Table 1.1: Newly acquired past history for finding 'credibility'

Remark 1.1.1

The shall leave the following as remarks.

- How is the policyholder's own experience account for? This is a key question that will be addressed in this course.
- Risks in a given rating class are not perfectly identical (i.e., no rating system is perfect)
- One may refine the rating system by incorporating more factors but it is time-consuming (and no system is perfect).

Thus, credibility theory is designed such that it

- accounts for heterogeneity within a given rating lass; and
- provides a theoretical justification to charge a premium that reflects to the policyholder's own experience.

1.2 Review of Probability

You are expected to be familiar with the following concepts:

- Joint and Marginal Distribution
- Conditional Distribution
- Mixture Distributions (see also ACTSC₄₃₁)
 - *n*-point Mixture
- Conditional Expectation

Some examples or more detailed review will be added for each topic if I come to work through them in detail.

Review of Statistics

In this chapter, we will review the following notions:

- Unbiased estimation
- Maximum likelihood estimation
- Bayesian estimation 🖈

2.1 Unbiased Estimation

Suppose we are given a **parametric model** ¹ of X, i.e. the distribution of $X \mid \Theta = \theta$ is known but θ is unknown. Furthermore, we have a **random sample** of X, i.e. we have $\{X_i\}_{i=1}^n$ is an independent and identically distributed (iid) sequence of random variables (rv) such that $X_i \sim X$.

¹ See ACTSC431.

■ Definition 1 (Estimate)

An estimate is a specific value that is obtained when applying an estimation procedure to a set of numbers, and in our case, rvs. We usually denote an estimate by a hat \cdot.

■ Definition 2 (Estimator)

An estimator is a rule or formula that produces an estimate. We usually denote an estimator by ~.

66 Note 2.1.1

An estimate is a number or a function, while an estimator is an rv or a random function.

Remark 2.1.1

In this course, we will not make a difference between the estimator and the estimate, and will use only .

■ Definition 3 (Biased and Unbiased Estimator)

We say that an estimator, $\hat{\theta}$, is unbiased if

$$E[\hat{\theta} \mid \theta] = \theta$$

for all θ . We say that an estimator is biased if it is not unbiased, and we define the bias as

$$bias_{\hat{\theta}}(\theta) = E[\hat{\theta} \mid \theta] - \theta.$$

Let's have ourselves a silly example.

Example 2.1.1

Let $(X_1, ..., X_n)$ be a random sample of $\text{Exp}(\beta)$. The sample mean

$$\overline{X} = \frac{1}{n} \sum_{i=1}^{n} X_i,$$

is an unbiased estimator for the mean β ; observe that by the **linearity** of the expectation, we have

$$E[\overline{X}] = E\left[\frac{1}{n}\sum_{i=1}^{n}X_{i}\right] = \frac{1}{n}\sum_{i=1}^{n}E[X_{i}] = \frac{1}{n}(n\beta) = \beta.$$

Example 2.1.2

Let $\{X_i\}_{i=1}^n$ be a random sample of $X \sim \text{Unif}(0, \theta)$. Let us construct two unbiased estimators for θ using

- 1. the sample mean \overline{X} ; and
- 2. order statistics $X_{(n)} \coloneqq \max_{1 \le i \le n} \{X_i\}$.



Solution

1. Observe that

$$E[\overline{X}] = E\left[\frac{1}{n}\sum_{i=1}^{n}X_i\right] = \frac{1}{n}\sum_{i=1}^{n}E[X_i] = \frac{1}{n}\cdot n\left(\frac{\theta}{2}\right) = \frac{\theta}{2}.$$

This tells us that if we picked $\hat{\theta} = 2\overline{X}$, then we would end up with

$$E[2\overline{X}] = \theta.$$

Thus $2\overline{X}$ is an unbiased estimator of θ .

2. Using the Darth Vader rule 2 , since the X_i 's form a random sample of X, and the bounds for each X_i is 0 and θ , we have that

$$\begin{split} E[X_{(n)}] &= \int_0^\infty \overline{F}_{X_{(n)}}(x) \, dx \\ &= \int_0^\infty \left(1 - P(\max\{X_1, X_2, \dots, X_n\}) \le x \right) dx \\ &= \int_0^\infty \left(1 - P(X_1 \le x) P(X_2 \le x) \dots P(X_n \le x) \right) dx \\ &= \int_0^\theta \left(1 - \left(\frac{x}{\theta} \right)^n \right) dx \\ &= \theta - \frac{1}{n+1} \left(\frac{x^{n+1}}{\theta^n} \right) \Big|_{x=0}^{x=\theta} = \frac{n}{n+1} \theta, \end{split}$$

where we note that we can change the bounds as such since $X \sim$ Unif $(0, \theta)$ implies that

$$P(X \le \theta) = \begin{cases} \frac{x}{\theta} & 0 \le x \le \theta \\ 1 & x > \theta \end{cases}.$$

Thus, to get an unbiased estimator for θ , we simply need to consider

$$\hat{\theta} = \frac{n+1}{n} X_{(n)},$$

which then

$$E\left[\frac{n+1}{n}X_{(n)}\right]=\theta.$$

² The **Darth Vader rule** is given as: if *X* is a non-negative rv, then

$$E[X] = \int_0^\infty \overline{F}_X(x) \, dx,$$

where \overline{F}_X is the survival function of X.

♦ Proposition 1 (Sample Mean as the Unbiased Estimator of the Mean)

Let $\{X_i\}_{i=1}^n$ be a random sample of X which has mean μ . Then \overline{X} is an unbiased estimator of μ .

Proof

We have that

$$E[\overline{X}] = \frac{1}{n} \sum_{i=1}^{n} E[X_i] = \frac{1}{n} (n\mu) = \mu.$$

■ Definition 4 (Sample Variance)

Let $\{X_i\}_{i=1}^n$ be a random sample of X which has mean μ and variance σ^2 . We define the sample variance as

$$\hat{\sigma}^2 := \frac{1}{n-1} \sum_{i=1}^n (X_i - \overline{X})^2.$$

♦ Proposition 2 (Sample Variance as the Unbiased Estimator of the Variance)

Let $\{X_i\}_{i=1}^n$ be a random sample of X which has mean μ and variance σ^2 . Then the sample variance $\hat{\sigma}^2$ is an unbiased estimator of σ^2 .

Proof

First, note that

$$Var(\overline{X}) = Var\left(\frac{1}{n}\sum_{i=1}^{n}X_{i}\right)$$
$$= \frac{1}{n^{2}}n Var(X_{i})$$
$$= \frac{1}{n}\sigma^{2}.$$

Thus

$$E\left[\sum_{i=1}^{n} (X_{i} - \overline{X})^{2}\right] = E\left[\sum_{i=1}^{n} (X_{i} - \mu + \mu - \overline{X})^{2}\right]$$

$$= \sum_{i=1}^{n} E\left[(X_{i} - \mu)^{2}\right] + \sum_{i=1}^{n} E\left[(\mu - \overline{X})^{2}\right]$$

$$+ 2E\left[\sum_{i=1}^{n} (X_{i} - \mu)(\mu - \overline{X})\right]$$

$$= n\sigma^{2} + n \operatorname{Var}(\overline{X})^{-3} + 2nE\left[(\overline{X} - \mu)(\mu - \overline{X})\right]^{-4}$$

$$= n\sigma^{2} - n \operatorname{Var}(\overline{X})$$

$$= n\sigma^{2} - n \left(\frac{1}{n}\sigma^{2}\right)$$

$$= (n-1)\sigma^{2}.$$

It follows that

$$E\left[\frac{1}{n-1}\sum_{i=1}^{n}(X_i-\overline{X})^2\right]=\sigma^2.$$

Remark 2.1.2

In general, unbiasedness is not preserved under parameter transformations. *E.g.*, $\frac{1}{\overline{X}}$ is generally not unbiased for μ , where μ is the mean of \overline{X} .

Some unbiased estimators can also be unreasonable.

Example 2.1.3

Consider $X \sim \text{Poi}(\lambda)$, where $\lambda > 0$. Note that

$$E[(-1)^X] = e^{\lambda(-1-1)} = e^{-2\lambda}$$

by the probability generating function method, and we see that $(-1)^X$ is an unbiased estimator of $e^{-2\lambda}$. However, we see that $(-1)^X$ only takes on values ± 1 , which is nowhere close to $e^{-2\lambda}$.

Intuitively, $e^{-2\overline{X}}$ would be a "better" estimator despite the fact that it is biased.

- ⁴ This relies on the fact that \overline{X} is the unbiased estimator for μ (cf.
- ♦ Proposition 1). We then use the definition of the variance to achieve
- ⁴ We used the fact that

$$\sum_{i=1}^{n} (X_i - \mu) = \sum_{i=1}^{n} X_i - n\mu = n\overline{X} - n\mu.$$

Also, note that

$$Var(\overline{X}) = E[(\overline{X} - \mu)^2].$$

Despite shortcomings like the above, unbiasedness is generally a good property for an estimator to have.

2.2 Mean Squared Error

■ Definition 5 (Mean Squared Error)

Suppose $\hat{\theta}$ is an estimator for the parameter θ . The mean squared error (MSE) of $\hat{\theta}$ is defined as

$$\mathrm{MSE}_{\hat{\theta}}(\theta) \coloneqq E\left\lceil (\hat{\theta} - \theta)^2 \right\rceil = \mathrm{Var}(\hat{\theta}) + \mathrm{bias}_{\hat{\theta}}(\theta)^2.$$

Proof

It is not immediately clear how the two expressions are the same. We shall prove it here. First, note that $\operatorname{bias}_{\hat{\theta}}(\theta) = E[\hat{\theta}] - \theta$ is a real value. Using a similar idea as in \bullet Proposition 2, we see that

$$E\left[\left(\hat{\theta} - \theta\right)^{2}\right] = E\left[\left(\hat{\theta} - E[\hat{\theta}] + E[\hat{\theta}] - \theta\right)^{2}\right]$$

$$= E\left[\left(\hat{\theta} - E[\hat{\theta}]\right)^{2}\right] + E\left[\left(E[\hat{\theta}] - \theta\right)^{2}\right]$$

$$+ 2E\left[\left(\hat{\theta} - E[\hat{\theta}]\right)\left(E[\hat{\theta}] - \theta\right)\right]$$

$$= Var(\hat{\theta}) + bias_{\hat{\theta}}(\theta)^{2}$$

$$+ 2bias_{\hat{\theta}}(\theta) E[\hat{\theta} - E[\hat{\theta}]]$$

$$= Var(\hat{\theta}) + bias_{\hat{\theta}}(\theta)^{2}.$$

66 Note 2.2.1

The MSE is a measure to evaluate the quality of estimators. The smaller the MSE, the better the estimator.

2.3 Maximum Likelihood Estimation

■ Definition 6 (Likelihood Function)

Let $\{X_i\}_{i=1}^n$ be a random sample of X with density $f(x;\underline{\theta})$, where $\underline{\theta}$ is possibly a vector of parameters. The likelihood function for $\underline{\theta}$ is defined as

 $L(\underline{\theta}) = \prod_{i=1}^{n} f(X_i; \underline{\theta}).$

■ Definition 7 (Maximum Likelihood Estimation)

The maximum likelihood estimation (MLE) of $\hat{\theta}$ of θ is an approach that maximizes $L(\hat{\theta})$.

66 Note 2.3.1

Heuristically, under the MLE, $\hat{ heta}$ is the most likely parameter for the sample (X_1, \ldots, X_n) to be realized.

Sometimes, the likelihood function is difficult to work with. Fortunately, since ln x is a increasing bijective function that preserves monotonicity, we can make us of this property to ensure maximality.

■ Definition 8 (Log-likelihood Function)

The log-likelihood function is defined as

$$l(\underline{\theta}) = \sum_{i=1}^{n} \ln(f(X_i; \underline{\theta})).$$

Example 2.3.1

Let $\{X_i\}_{i=1}^n$ be a random sample for $N(\mu, v)$. Find the MLE for $\mu, v.$

Solution

First, we shall work on getting an MLE for μ . The likelihood function here is

$$L(\mu) = \prod_{i=1}^{n} f(X_i; \mu)$$

$$= \prod_{i=1}^{n} \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(X_i - \mu)^2}{2\sigma^2}}$$

$$\propto \exp\left\{-\frac{1}{2\sigma^2} \sum_{i=1}^{n} (X_i - \mu)^2\right\}.$$

Evaluating the derivative and equating it to 0 would be fruitless, since this is an exponentiation. Thus we appeal to the log-likelihood, which is

$$l(\mu) \propto \sum_{i=1}^{n} (X_i - \mu)^2.$$

The derivative log-likelihood is thus

$$\frac{dl}{d\mu} \propto -2\sum_{i=1}^{n} (X_i - \mu).$$

Equating the above to 0, we get

$$\hat{\mu} = \overline{X}$$
.

Now for an MLE of σ^2 . For sanity, let us denote $\tau = \sigma^2$. Then the likelihood function, focusing on τ , is

$$L(\tau) = \prod_{i=1}^{n} \frac{1}{\sqrt{2\pi\tau}} e^{-\frac{(X_i - \mu)^2}{2\tau}}$$
$$\propto \tau^{-\frac{n}{2}} e^{-\frac{1}{2\tau} \sum_{i=1}^{n} (X_i - \mu)^2}.$$

Again, the likelihood involves an exponentiation, so we appeal to the log-likelihood, which is

$$l(\tau) \propto -\frac{n}{2} \ln \tau - \frac{1}{2\tau} \sum_{i=1}^{n} (X_i - \mu)^2.$$

The derivative of the log-likelihood is

$$\frac{dl}{d\tau} = -\frac{n}{2\tau} + \frac{1}{2\tau^2} \sum_{i=1}^{n} (X_i - \mu)^2.$$

Equating the above to 0, we get

$$n = \frac{1}{\hat{\tau}} \sum_{i=1}^{n} (X_i - \hat{\mu})^2,$$

and so

$$\hat{\sigma}^2 = \hat{\tau} = \frac{1}{n} \sum_{i=1}^n (X_i - \overline{X})^2$$

Bayesian Estimation

From Klugman et al. 2012,

The Bayesian approach assumes that only the data actually observed are relevant and it is the population distribution that is variable.

■ Definition 9 (Prior Distribution)

The prior distribution is a probability distribution over the space of possible parameter values. It is denoted $\pi(\theta)$ and represents our opinion concerning the relative chances that various values of θ are the true value of the parameter.

66 Note 2.4.1

- The parameter θ may be scalar or vector valued.
- Determining the prior distribution has always been one of the barriers to the widespread acceptance of the Bayesian methods, since it is almost certainly the case that your experience has provided you with some insight about possible parameter values before the first data point has been observed.

We shall use the following concepts from multivariate statistics to obtain the following definitions.

■ Definition 10 (Joint Distribution)

Let $\{X_i\}_{i=1}^n$ be a random sample of the $rv\ X$, and Θ another $rv\ that$ is independent of the X_i 's 5 , with pdf π . Let $\vec{X}=(X_1,X_2,\ldots,X_n)$. Then the joint distribution of \vec{X} and Θ is defined as

$$f_{\vec{X},\Theta}(\vec{x},\theta) = f_{\vec{X}\mid\Theta}(\vec{x}\mid\theta)\pi(\theta).$$

⁵ Note that Θ does not necessarily have a similar distribution to X.

■ Definition 11 (Marginal Distribution)

Let $\{X_i\}_{i=1}^n$ be a random sample of the $rv\ X$, and Θ another $rv\ that$ is independent of the X_i 's 6 , with pdf π . Let $\vec{X}=(X_1,X_2,\ldots,X_n)$. Then the marginal distribution of \vec{X} is defined as

$$f_{\vec{X}}(\vec{x}) = \int f_{\vec{X}|\Theta}(\vec{x} \mid \theta) \pi(\theta) d\theta.$$

 6 Note that Θ does not necessarily have a similar distribution to X.

Once we have obtained data, we can look back at our prior distribution and "update" it to...

■ Definition 12 (Posterior Distribution)

Let $\{X_i\}_{i=1}^n$ be a random sample of the rv X, and Θ another rv that is independent of the X_i 's 7 , with pdf π . The posterior distribution, denoted by $\pi_{\Theta \mid \vec{X}}(\theta \mid \vec{x})$, is the conditional probability distribution of the parameters given the observed data.

 7 Note that Θ does not necessarily have a similar distribution to X.

It is easy to find out what the general formula of the posterior distribution is. One simply needs to make use of E Definition 10 and Definition 11. The proof of the following proposition is left as an easy brain exercise for the reader.

Exercise 2.4.1

Prove Proposition 3.

• Proposition 3 (Formula for the Posterior Distribution)

With the assumptions in 🗏 Definition 12, we have that the posterior distribution can be computed as

$$\begin{split} \pi_{\Theta \mid \vec{X}}(\theta \mid \vec{x}) &= \frac{f_{\vec{X},\Theta}(\vec{x},\theta)}{f_{\vec{X}}(\vec{x})} \\ &= \frac{\left(\prod_{i=1}^{n} f_{X_{i}\mid\Theta}(x_{i}\mid\theta)\right)\pi(\theta)}{\int_{\forall \theta} \left(\prod_{i=1}^{n} f_{X_{i}\mid\Theta}(x_{i}\mid\theta)\right)\pi(\theta)\,d\theta}. \end{split}$$

■ Definition 13 (Posterior Mean)

The posterior mean is defined as the expected value of the posterior distribution.

■ Definition 14 (Bayes Estimator)

The Bayes estimator of Θ is the posterior mean of Θ , defined as

$$\hat{\theta}_{B} := E[\Theta \mid \vec{X} = \vec{x}] = \int_{\forall \theta} \theta \cdot \pi_{\Theta \mid \vec{X}}(\theta \mid \vec{x}).$$

66 Note 2.4.2

It can be shown that $\hat{\theta}_B$ minimizes the mean square error

$$\min_{\hat{\theta}} E\left[\left(\hat{\theta} - \Theta\right)^2 \mid \vec{X} = \vec{x}\right].$$

2.4.1 Conjugate Prior Distributions and the Linear Exponential Family

■ Definition 15 (Conjugate Prior Distribution)

A prior distribution is said to be a conjugate prior distribution for a given model if the resulting posterior distribution is from the same family as the prior, although possibly with different parameters.

8

⁸ More examples should be added here.

Example 2.4.1

The following are some important/prominent examples of conjugate prior distributions:

$\pi(\theta)$	$f(x \mid \theta)$	$\pi(\theta \mid \vec{x})$
Gamma	Poisson	Gamma
Normal	Normal	Normal
Beta	Binomial	Beta
Beta	Geometric	Beta

Table 2.1: Important/Prominent Conjugate Prior Distributions

■ Definition 16 (Linear Exponential Family)

An rv X is said to belong to the linear exponential family if its pdf is of the form

$$f(x,\theta) = \frac{p(x)e^{xr(\theta)}}{q(\theta)},$$

where p(x) is some function of x, and $r(\theta)$, $q(\theta)$ are some functions of θ , and the support of f does not depend on θ .

Example 2.4.2

Some members of the linear exponential family include

- $\operatorname{Exp}(\theta): f(x, \theta) = \frac{1}{\theta} e^{-\frac{x}{\theta}}$, where p(x) = 1, $r(\theta) = -\frac{1}{\theta}$ and $q(\theta) = \theta$.
- $Gam(\alpha, \theta)$: $f(x, \alpha, \theta) = \frac{1}{\Gamma(\alpha)\theta^{\alpha}} x^{\alpha 1} e^{-\frac{x}{\theta}}$.
- $\operatorname{Poi}(\theta) : f(x, \theta) = \frac{\theta^x e^{-\theta}}{x!} = \frac{\frac{1}{x!} e^{x \ln \theta}}{e^{\theta}}$
- $N(\theta, v) : f(x, \theta, v) = \frac{1}{\sqrt{2\pi v}} e^{-\frac{(x-\theta)^2}{2v}} = \frac{(2\pi v)^{-\frac{1}{2}} e^{-\frac{x^2}{2v}} e^{x\frac{\theta}{v}}}{e^{\theta^2/2v}}$

66 Note 2.4.3

Basically, functions the belong to a linear exponential family is a linear-like function with an exponent.

Theorem 4 (Conjugate Prior Distributions of Linear Exponential Distributions)

Suppose that given $\Theta= heta$ the rvs $ec{X}$ are iid with pf

$$f_{X_j \mid \Theta}(x_j \mid \theta) = \frac{p(x_j)e^{r(\theta)x_j}}{q(\theta)},$$

where Θ has the pdf

$$\pi(\theta) = \frac{[q(\theta)]^{-k} e^{\mu k r(\theta)} r'(\theta)}{c(\mu, k)},$$

where μ and k are parameters of the distribution and $c(\mu, k)$ is the nor*malizing constant* ⁹. Then the posterior pf $\pi_{\Theta \mid \vec{X}}(\theta \mid \vec{x})$ is of the same form as $\pi(\theta)$, i.e. $\pi(\theta)$ is a conjugate prior distribution function.

⁹ The normalizing constant is used to reduce any probability function to a probability density function with a total probability of 1. (Source: Wikipedia)

Proof

Notice that the posterior distribution is

$$\pi(\theta \mid \vec{x}) = \frac{\left(\prod_{i=1}^{n} f_{X_{i}\mid\Theta}(x_{i}\mid\theta)\right)\pi(\theta)}{\int_{\forall\theta} \left(\prod_{i=1}^{n} f_{X_{i}\mid\Theta}(x_{i}\mid\theta)\right)\pi(\theta)d\theta}$$

$$\propto \left(\prod_{i=1}^{n} f_{X_{i}\mid\Theta}(x_{i}\mid\theta)\pi(\theta)\right)$$

$$= \left(\prod_{i=1}^{n} \frac{p(x_{j})e^{r(\theta)x_{j}}}{q(\theta)}\right) \left(\frac{[q(\theta)]^{-k}e^{\mu k r(\theta)}r'(\theta)}{c(\mu,k)}\right)$$

$$\propto q(\theta)^{-(n+k)}e^{\mu k + n\bar{x}r(\theta)}r'(\theta)$$

$$= q(\theta)^{-k^{*}}e^{\mu^{*}k^{*}r(\theta)}r'(\theta),$$

where

$$k^* = k + n$$
, and $\mu^* = \frac{\mu k + \sum x_j}{k + n} = \frac{k}{k + n} \mu + \frac{n}{k + n} \overline{x}$,

and we see that the posterior distribution has the same form as $\pi(\theta)$.

Example 2.4.3

One non-example is mentioned in Example 2.4.1: the distribution of X_i is not from the linear exponential family, but we still obtain that the posterior distribution has a similar distribution to the posterior distribution.

Limited Fluctuation Credibility The-

ory

The Limited Fluctuation Credibility Theory provides a mechanism for assigning **full** or **partial credibility** to a policyholder's experience. The difficulty with this approach is its lack of a sound underlying mathematical theory that justifies the use of these methods. Despite that fact, it is still widely used today, especially in the United States.

3.1 Limited Fluctuation Credibility

From Klugman et al. 2012,

This branch of credibility theory represents the first attempt to quantify the credibility problem.

This approach is also known as the "American credibility". It was first proposed by Mowbray in 1914 ¹.

The problem can be formulated as follows. Suppose that $\{X_i\}_{i=1}^n$ represents a policyholder's claim amounts in the past n years. Furthermore, we assume that the X_i 's have

- the same expected value, i.e. $E[X_i] = \mu$ for some μ ; and
- variance, i.e. $Var(X_i) = \sigma^2$ for some σ .

From our revision in the last section, we know that \overline{X} is an unbiased estimator for μ , and if the X_i 's are independent, then $Var(\overline{X}) = \frac{\sigma^2}{n}$.

The goal here is to figure our how much to charge for the next

¹ Mowbray, A. H. (1914). How extensive a payroll exposure is necessary to give a dependable pure premium? *Proceedings* of the Casualty Actuarial Society, I:24–30 premium, i.e. determining $E[X_{n+1}]$. We have at least the following 3 possibilities:

- ignore past data (no credibility) and charge M, a value, called the manual premium ², obtained from experience on other similar but non-identical policyholders;
- ignore M and charge \overline{X} (full credibility); and a third possibility is to
- choose some combination of M and \overline{X} (partial credibility).

From the POV of an insurer, it seems sensible to favor \overline{X} if the experience is "stable", i.e. there is little fluctuation, represented by a small σ^2 . Stable values imply that \overline{X} is more reliable as a predictor. Conversely, if \overline{X} is volatile, then M would be a safer choice.

² This name is obtained from the fact that it usually comes from a book (manual) of premiums.

3.2 Full Credibility

In full credibility theory, there are only 2 outcomes: either we

- assign full credibility, that is to charge \overline{X} ; or
- no credibility, where we charge M.

One method to 'quantify the stability' of \overline{X} 3 is to infer that \overline{X} is stable if the difference between \overline{X} and μ is small relative to μ with high probability, i.e.

$$P(\left|\overline{X} - \mu\right| \le \varepsilon \mu) \ge p \tag{3.1}$$

for some $\varepsilon > 0$ and 0 . We may rewrite Equation (3.1) as

$$P\left(\frac{\left|\overline{X}-\mu\right|}{\sigma/\sqrt{n}} \le \frac{\varepsilon\mu}{\sigma/\sqrt{n}}\right) \ge p.$$

Now let y_p be defined as by

$$y_p = \operatorname{VaR}_p\left(\frac{\left|\overline{X} - \mu\right|}{\sigma/\sqrt{n}}\right) = \inf\left\{y \in \mathbb{R} : P\left(\frac{\left|\overline{X} - \mu\right|}{\sigma/\sqrt{n}} \le y\right) \ge p\right\}.$$

If \overline{X} is continuous, then the \geq sign above can be replaced with an "=" sign ⁴, and y_p satisfies

⁴ See ACTSC431.

³ This has become the standard method for 'quantifying stability' for \overline{X} .

$$P\left(\frac{\left|\overline{X}-\mu\right|}{\sigma/\sqrt{n}} \le y_p\right) = p. \tag{3.2}$$

Then the condition for full credibility is

$$y_p \leq \frac{\varepsilon \mu}{\sigma/\sqrt{n}}$$
.

Making n the subject, we have that the number of exposure required for full credibility is thus

$$n \ge \left(\frac{y_p}{\varepsilon}\right)^2 \frac{\sigma^2}{\mu^2} = \lambda_0 \frac{\sigma^2}{\mu^2},\tag{3.3}$$

where we let $\lambda_0 = \left(\frac{y_p}{\varepsilon}\right)^2$ for notational succinctness since it is a constant that depends only p and ε .

It is often difficult to identify a distribution for \overline{X} , of which y_p depends on. Recall the normal approximation, which is applicable if *n* is large ⁵:

$$\frac{\overline{X} - \mu}{\sigma / \sqrt{n}} \approx Z_{0,1} \sim N(0,1)$$

Then Equation (3.2) becomes

$$p = P(|Z| \le y_p) = \Phi(y_p) - \Phi(-y_p)$$

= $\Phi(y_p) - 1 + \Phi(y_p) = 2\Phi(y_p) - 1$.

Thus

$$y_p pprox \Phi^{-1}\left(\frac{1+p}{2}\right).$$

Example 3.2.1

Suppose that one has data $\{X_i\}_{i=1}^{10}$ on the claim amounts in the last 10 periods, where

$$X_i = 0 \text{ for } i = 1, \dots, 6,$$

and

$$X_7 = 253$$
, $X_8 = 398$, $X_9 = 439$, $X_{10} = 756$.

Determine the condition for full credibility with $\varepsilon = 0.05$ and p =0.9.

Solution

P (Condition for Full Credibility)

- 1. Use the central limit theorem argument
- 2. Calculate RHS of Equation (3.3).

5 Is this not circular!?

We need to first determine the sample mean and sample variance, and we shall use the unbiased estimators of μ and σ^2 respectively: they are

$$\overline{X} = \frac{1}{10} \sum_{i=1}^{10} X_i = \frac{0 + 253 + 398 + 439 + 756}{10} = 184.6,$$

and

$$\hat{\sigma}^2 = \frac{1}{10-1} \sum_{i=1}^{10} (X_i - \overline{X})^2 = 267.89^2.$$

We also need

$$y_p = \Phi^{-1}\left(\frac{1+p}{2}\right) = \Phi^{-1}(.95) = 1.645.$$

Then we require that

$$n \ge \left(\frac{1.645}{0.05}\right)^2 \left(\frac{267.89^2}{184.6^2}\right) = 2279.5.$$

We see that the 10 observations definitely do not deserve full credibility.

Full credibility is sometimes given on a number of claims basis (instead of on the claims amount).

Example 3.2.2

Suppose that one has iid data $\{N_i\}_{i=1}^n$ on the number of claims in the past n periods, with $N_i \sim \operatorname{Poi}(\lambda)$. Determine the condition for full credibility in terms of the **expected total number of claims** given that p = 0.9 and $\varepsilon = 0.05$.

Solution

Since $N_i \sim \text{Poi}(\lambda)$, we have $E[N_i] = \lambda = \text{Var}(N_i)$. Furthermore,

$$y_p = \Phi^{-1}(0.95) = 1.645.$$

Now since the condition is

$$n \ge \lambda_0 \frac{\sigma^2}{u^2} = \frac{\lambda_0}{\lambda},$$

and we want the expected total number of claims, we focus on look-

ing at

$$n\mu = n\lambda \ge \lambda_0$$
.

Observe that

$$\lambda_0 = \left(\frac{1.645}{0.05}\right)^2 = 1082.41,$$

we have that the required expected total number of claims should fulfill

$$n\lambda \ge 1082.41$$
.

Example 3.2.3 (Compound Poisson for Full Credibility)

Let $\{X_i\}_{i=1}^n$ be a sequence of iid compound Poisson rvs, given by

$$X_i = \sum_{j=1}^{N_i} Y_{i,j} = \begin{cases} \sum_{j=1}^{N_i} Y_{i,j}, & N_i \ge 0 \\ 0 & N_i = 0 \end{cases}$$

where

- $\{N_i\}_{i=1}^n$ are iid with $N_i \sim \text{Poi}(\lambda)$ for each i; and
- $\{Y_{i,j}\}$ are also iid with mean μ_Y and variance σ_Y^2 .

Determine the condition for full credibility.

Solution

We require the unconditional sample mean and sample variance of X_i ; they are

$$E[X_i] = E[E[X_i \mid N_i]] = E[N_i]E[Y_{i,j}] = \lambda \mu_Y,$$

and

$$Var(X_i) = Var(E[X_i \mid N_i]) + E[Var(X_i \mid N_i)]$$

$$= Var(N_i\mu_Y) + E[N_i\sigma_Y^2]$$

$$= \mu_Y^2\lambda + \sigma_Y^2\lambda$$

$$= \lambda(\mu_Y^2 + \sigma_Y^2).$$

Thus, the condition for full credibility is

$$n \geq \lambda_0 rac{\lambda(\mu_Y^2 + \sigma_Y^2)}{\lambda^2 \mu_Y^2} = rac{\lambda_0}{\lambda} \left(1 + rac{\sigma_Y^2}{\mu_Y^2}
ight).$$

To further illustrate that we can use the concept of full credibility for different things, the following example is provided.

Example 3.2.4

Suppose that the average claim size for a group of insureds is 1500 with a standard deviation of 7500. Furthermore, assume that claim counts have a Poisson distribution. For $\varepsilon=0.06$ and p=0.9, determine the standard for full credibility based on the

- 1. total claim amount; and
- 2. total number of claims,

in terms of the expected total number of claims.



1. Using the last example and letting

$$E[X_i] = \mu$$
 and $Var(X_i) = Var(X_i) = \sigma^2$,

the standard for full credibility is

$$n \ge \frac{\lambda_0}{\lambda} \left(1 + \frac{\sigma_Y^2}{\mu_Y^2} \right).$$

We are given that

$$\mu_Y = 1500$$
 and $\sigma_Y^2 = 7500^2$.

Thus

$$n \ge \frac{1.645^2}{0.06^2 \lambda} \left(1 + \frac{7500^2}{1500^2} \right) = \frac{19543.51}{\lambda}.$$

In terms of the expected total number of claims, we have

$$n\lambda > 19543.51$$
.

Thus the observed total number of claims of past claims must be at least 19544 to assign full credibility.

2. Using Example 3.2.2, we have

$$n \ge \frac{\lambda_0}{\lambda} = \frac{751.67}{\lambda}.$$

Thus, in terms of the expected total number of claims, we have

$$n\lambda \ge 751.67$$
.

Therefore, the observed total number of past claims must be at least 752 to assign full credibility.

3.3 Partial Credibility

If full credibility is inappropriate, then we may want to assign partial **credibility** to the past experience \overline{X} in the net premium. Without much mathematical support, it was suggested that we let the net premium be defined as a weighted average of \overline{X} and the manual premium M, i.e.

$$P = Z\overline{X} + (1 - Z)M,$$

where $Z \in [0,1]$ is known as the credibility factor ⁶ ⁷, which is a value that needs to be chosen.

In the actuarial literature ⁸, there are various suggestions for determining Z. However, they are usually justified on intuition rather than theoretically sound grounds. We shall discuss one of the choices here, which is flawed, but is at least simple.

Recall that the goal of the full-credibility standard is to ensure that the difference between \overline{X} and μ is small with high probability (cf. beginning of Section 3.2). Since \overline{X} is unbiased, to achieve this standard is basically ⁹ equivalent to controlling the variance of \overline{X} . Note that full credibility fails when

$$n < \lambda_0 \left(\frac{\sigma^2}{\mu^2} \right), \tag{3.4}$$

and since the sample variance (which is unbiased for the variance) is

$$\operatorname{Var}(\overline{X}) = \frac{\sigma^2}{n},$$

rearranging Equation (3.4), we have that

$$\operatorname{Var}(\overline{X}) = \frac{\sigma^2}{n} > \frac{\mu^2}{\lambda_0}.$$

⁶ It is important to note there that *Z* is not an rv. It is simply a pretentious choice of notation for what is to come. 7 It is interesting to remark that Mowbray 1914 considered full but not partial credibility.

⁸ Klugman, S. A., Panjer, H. H., and Willmot, G. E. (2012). Loss Models: From Data to Decisions. John Wiley & Sons Inc., Hoboken, New Jersey, 4th edition

⁹ This is exactly the case if \overline{X} is normal.

Thus, we choose *Z* such that it controls the variance of the credibility premium as such:

$$\frac{\mu^2}{\lambda_0} = \text{Var}(P) = \text{Var}(Z\overline{X} + (1 - Z)M)$$
$$= Z^2 \text{Var}(\overline{X}) = Z^2 \cdot \frac{\sigma^2}{n}.$$

Thus, since we want Z as a weighted average, we let

$$Z = \min \left\{ \frac{\mu}{\sigma} \sqrt{\frac{n}{\lambda_0}}, 1 \right\}.$$

10 Note that

$$\frac{\mu}{\sigma}\sqrt{\frac{n}{\lambda_0}} = \sqrt{\frac{n}{\lambda_0(\frac{\sigma^2}{\mu^2})}},$$

which is the square root of the actual number of exposures divided by the number of exposures needed for full credibility. This is also referred to as the Square-root rule for partial credibility.

Example 3.3.1

Suppose that past observations of the number of claims $\{N_i\}_{i=1}^n$ are iid and $N_i \sim \text{Poi}(\lambda)$. Determine the credibility factor Z based on the number of claims.

Solution

Note that

$$\mu = E[N_i] = \lambda$$
 and $\sigma^2 = Var(N_i) = \lambda$.

We have that

$$Z = \min \left\{ \frac{\mu}{\sigma} \sqrt{\frac{n}{\lambda_0}}, 1 \right\} = \min \left\{ \sqrt{\frac{n\lambda}{\lambda_0}}, 1 \right\}.$$

Example 3.3.2

Consider the setup in Example 3.2.3. Determine the credibility factor Z based on the amount of claims.

Solution

¹⁰ Note that this choice of *Z* has some consistency with full credibility, since Z=1 iff $n \geq \lambda_0 \frac{\sigma^2}{\sigma^2}$.

We have that

$$\mu = E[X_i] = \lambda \mu_Y$$
 and $\sigma^2 = Var(X_i) = \lambda(\mu_Y^2 + \sigma_Y^2)$.

Then since

$$rac{\mu}{\sigma}\sqrt{rac{n}{\lambda_0}}=\sqrt{rac{n\lambda}{\lambda_0}\cdotrac{\mu_Y^2}{\mu_Y^2+\sigma_Y^2}},$$

we have that

$$Z = \min \left\{ \sqrt{rac{n\lambda}{\lambda_0} \cdot rac{\mu_Y^2}{\mu_Y^2 + \sigma_Y^2}}, 1
ight\}$$

DIFFERENT CREDIBILITY FACTORS may arise depending on the basis of which the credibility is founded upon.

Example 3.3.3

Consider the setup in Example 3.2.4. Further suppose that

- in thelast year, this group of insureds had 600 claims and a total loss of 15600; and
- the prior estimate of the total loss was 16500 (this is *M*).

Estimate the credibility premium for the group based on the

- 1. total claim amount; and
- 2. total number of claims.

Solution

1. We are given that $\mu_Y = 1500$, $\sigma_Y = 7500$ and $n\lambda = 600$. Thus

$$Z = \min \left\{ \sqrt{\frac{n\lambda}{\lambda_0} \cdot \frac{\mu_Y^2}{\mu_Y^2 + \sigma_Y^2}}, 1 \right\}$$
$$= \min \left\{ \sqrt{\frac{600}{\left(\frac{1.645}{0.06}\right)^2} \cdot \frac{1500^2}{1500^2 + 7500^2}}, 1 \right\}$$
$$= 0.17522$$

Thus the credibility premium for the group is

$$P = 0.17522\overline{X} + (1 - 0.17522)M$$

$$= 0.17522(15600) + (1 - 0.17522)(16500)$$
$$= 16342.302$$

11

2. Based on the total number of claims, the credibility factor is

$$Z = \min\left\{\sqrt{\frac{n\lambda}{\lambda_0}}, 1\right\} = \min\left\{\sqrt{\frac{600}{\left(\frac{1.645}{0.06}\right)^2}}, 1\right\} = 0.89343.$$

Thus the credibility premium for the group is

$$P = 0.89343\overline{X} + (1 - 0.89343)M = 15696.$$

3.4 Problems with Limited Fluctuation Credibility

 There is no theoretical model for the distribution of X_i's, and so there is no reason why

$$P = Z\overline{X} + (1 - Z)M$$

is a reasonable and more preferable to *M*.

- The choice of *Z* is rather arbitrary.
- There is no guidance to the choices of ε and p.
- The limited fluctuation approach does not examine the difference between μ and M. Furthermore, it is usually the case that M is also an estimate, and hence unreliable in itself.

 $\overline{X}=15600$ in this case, since this is the total loss over '1' period of time, in particular it is the total amount up to the latest time.

Bibliography

Klugman, S. A., Panjer, H. H., and Willmot, G. E. (2012). *Loss Models:* From Data to Decisions. John Wiley & Sons Inc., Hoboken, New Jersey, 4th edition.

Mowbray, A. H. (1914). How extensive a payroll exposure is necessary to give a dependable pure premium? *Proceedings of the Casualty Actuarial Society*, I:24–30.

Index

American credibility, 27

normalizing constant, 25

Bayes Estimator, 23

Bias, 14

Biased Estimator, 14

Posterior Mean, 23

pure premium, 9

Prior Distribution, 21

partial credibility, 33

Posterior Distribution, 22

Conjugate Prior Distribution, 24

credibility factor, 33

Credibility Theory, 9

rating class, 10

Darth Vader rule, 15

Sample Variance, 16

Square-root rule for partial credi-

Estimate, 13 bility, 34

Estimator, 13

20011110001, 1)

full credibility, 28

Unbiased Estimator, 14

Joint Distribution, 22

Likelihood Function, 19

Linear Exponential Family, 24

Log-likelihood Function, 19

manual premium, 10, 28

Marginal Distribution, 22

Maximum Likelihood Estimation,

19

Mean Squared Error, 18